

Chapter 1: How SourceMind Works

SourceMind turns a PDF into a workbook in three steps: upload, outline, read. You upload a PDF, and within seconds you have a full table of contents and every chapter's original text ready to read - no waiting, no AI calls, no cost. This document you're reading right now was ingested exactly the same way.

That speed is not an accident. Ingest is deterministic: SourceMind never asks an AI model to read your document just to figure out its structure. If your PDF has bookmarks (a table of contents embedded by whoever made the PDF), SourceMind uses those directly as chapter boundaries. If it doesn't, SourceMind falls back to even page windows - a fixed number of pages per chapter - so you always get a usable outline immediately, even from a raw scan with no bookmarks at all.

Either way, the text you see in each chapter is the real, original text extracted from your PDF - never rewritten, summarized, or paraphrased by a model. SourceMind calls this body text immutable: once ingested, it never changes unless you re-upload or re-ingest the source. Anything generated later - a lesson, flashcards, a quiz - is stored separately, so your source material is always there, unedited, to check against.

If a chapter title looks wrong, or a chapter split lands in an odd place, you can fix it. The outline supports renaming, reordering, merging adjacent chapters, and splitting one chapter into two - directly from the reader, without re-uploading anything.

Why go to this much trouble to avoid using AI for something AI is arguably good at? Cost and reproducibility. Ingesting a 500-page textbook this way costs nothing and takes seconds, whether you do it once or a hundred times. Every dollar SourceMind spends on an AI model is a deliberate choice you make later - never a side effect of simply opening a file.

That's it for ingest. The next chapter covers what SourceMind can generate for you once you decide you want it - lessons, flashcards, and quizzes - and how it keeps that spending visible and bounded.

Chapter 2: Lessons, Flashcards and Quizzes

Everything in this chapter is optional. SourceMind never calls an AI model on your behalf without you asking for it, chapter by chapter. Reading the raw source text costs nothing and always works - generation is there for when you want more: a plain-language lesson, a set of review flashcards, or a quiz to test yourself.

Before you generate anything, SourceMind can estimate what it will cost, based on the length of the chapter and the model you've configured. Once you do generate something, the exact token usage and estimated cost for that call is recorded and visible - nothing is hidden after the fact.

If you've set a spend cap for a course, SourceMind checks it immediately before every generation call and stops issuing new ones once the cap is reached, rather than letting a batch job run past a limit you set. A cap is a safety net, not a perfectly exact billing meter - but it means a runaway 'generate everything' click can't silently blow past what you were willing to spend.

Flashcards work with spaced repetition: after you generate a set for a chapter, you review them and grade your recall on a simple four-point scale - Again, Hard, Good, or Easy. SourceMind uses that grade to schedule when the card comes back: cards you find hard resurface sooner, cards you know well drift further out.

Quizzes are multiple-choice, scored automatically the moment you submit - answers are compared deterministically, so your score never depends on interpretation. You can retake a quiz on the same material as many times as you like.

One more thing worth knowing: if you regenerate a chapter's flashcards later - after editing the outline, or just wanting a fresh set - any card whose question and answer come out identical to before keeps its review history. Only cards that actually changed start over. Your progress is never thrown away by accident.

Chapter 3: Reading Efficiently

SourceMind remembers where you left off. As you read, your position is saved automatically in the background - per chapter, as a fraction of how far down the page you've scrolled, not a raw pixel count, so it resumes correctly no matter your window size or zoom level. Close the tab, come back tomorrow, and you're back exactly where you stopped.

The review queue works the same way. It's ordered by what's due, then by when each card was created - a fixed, stable order that doesn't reshuffle between visits. That means you can review ten cards, close the app, and come back later to pick up exactly where you left off in the queue, with no extra bookmarking required.

SourceMind is built keyboard-first - press ? anywhere to open a shortcuts overlay for whatever screen you're on. In the reader: the left/right arrows (or j and k) move between chapters, s switches between the source text and its generated lesson, and c opens or closes the course chat. In a review session: Space reveals a card's answer, then 1 through 4 grade your recall (Again, Hard, Good, Easy). In a quiz: 1 through 4 pick an answer choice, and Enter moves to the next question - or submits the quiz on the last one. Every control also remains reachable with plain Tab and Enter if you prefer.

That's the whole loop: upload a PDF, read it for free the moment it's ingested, and generate the extras - lessons, flashcards, quizzes - only when and where you actually want them, with your spending visible the whole way. This sample course will still be here under My Courses if you want to come back to it; deleting it won't bring it back automatically, so feel free to explore without worrying about tidying up afterward.